

AWS
re:Invent

O P N 3 0 3

BPF performance analysis at Netflix

Brendan Gregg

Senior Performance Architect
Netflix

Superpowers Demo

Agenda

Why BPF is changing linux


BPF internals

Performance analysis


Tool development

Why BPF is changing linux


50 years, one (dominant) OS model




Origins: Multics 1960s




Modern Linux: a new OS model



50 years, one process state model



BPF uses a new program state model



Netconf 2018

Alexei Starvoitov



Touch here to start presentation

BPF verifier in the future

- move away from existing brute force "walk all instructions" approach
- technology and static analysis
- remove `#define BPF_COMPLEXITY_LIMIT 128k` crutch
- remove `#define BPF_MAXINSNS 4k`
- support arbitrary large programs and libraries
 - 1 Million BPF instructions
- an algorithm to solve Rubik's cube will be expressible in BPF



BPF at Facebook

- ~40 BPF programs active on every server.
- ~100 BPF programs loaded on demand for short period of time.
- Mainly used by daemons that run on every server.
- Many teams are writing and deploying them.



Schedu

ftrace: Where modifying a running kerri

Analyzing changes to the binary interfa

BPF at Facebook - Alexei Starovoitov



Kernel Recipes 2019, Alexei Starovoitov

~40 active BPF programs on every Facebook server



>150K Amazon EC2 server instances


~34% US Internet traffic at night

>130M subscribers


~14 active BPF programs on every instance (so far)



Modern Linux: Event-based Applications



Modern Linux is becoming microkernel-ish



The word “microkernel” has already been invoked by Jonathan Corbet, Thomas Graf, Greg Kroah-Hartman, ...




BPF internals

BPF 1992: Berkeley Packet Filter

```
# tcpdump -d host 127.0.0.1 and port 80
(000) ldh      [12]
(001) jeq      #0x800          jt 2    jf 18
(002) ld      [26]
(003) jeq      #0x7f000001    jt 6    jf 4
(004) ld      [30]
(005) jeq      #0x7f000001    jt 6    jf 18
(006) ldb      [23]
(007) jeq      #0x84           jt 10   jf 8
(008) jeq      #0x6            jt 10   jf 9
(009) jeq      #0x11           jt 10   jf 18
(010) ldh      [20]
(011) jset     #0xffff         jt 18   jf 12
(012) ldxb    4*([14]&0xf)
(013) ldh      [x + 14]
(014) jeq      #0x50           jt 17   jf 15
(015) ldh      [x + 16]
(016) jeq      #0x50           jt 17   jf 18
(017) ret     #262144
(018) ret     #0
```


A limited
virtual machine for
efficient packet filters

BPF 2019: aka extended BPF



& Facebook Katran, Google KRSI, Netflix flowsrus,
and many more


BPF 2019



BPF is **open source** and in the **Linux kernel**
(you're all getting it)

BPF is also now a technology name,
and no longer an acronym

BPF Internals



A large, dark pile of coins is shaped like a question mark. The coins are stacked in a way that creates a textured, layered appearance, with some coins visible on top and others hidden in the shadows of the curve.

Is BPF Turing complete?



BPF: a new type of software

	Execution model	User-defined	Compile	Security	Failure mode	Resource access
User	task	yes	any	user-based	abort	syscall, fault
Kernel	task	no	static	none	panic	direct
BPF	event	yes	JIT, CO-RE	verified, JIT	error message	restricted helpers

Performance analysis


BPF enables a new class of
custom, efficient, and production-safe
performance analysis tools

BPF Performance Tools



Tool examples by subsystem

1. CPUs (scheduling)
2. Memory
3. Disks
4. File systems
5. Networking
6. Languages
7. Applications
8. Kernel
9. Hypervisors
10. Containers



Tool extensions & sources

.py: BCC (Python)

.bt: bpftrace

(some tools exist for both)

<https://github.com/iovisor/bcc>

<https://github.com/iovisor/bpftrace>

<https://github.com/brendangregg/bpf-perf-tools-book>

CPUs: execsnoop

New process trace

```
# execsnoop.py -T
TIME(s) PCOMM          PID  PPID  RET  ARGS
0.506  run            8745  1828  0   ./run
0.507  bash           8745  1828  0   /bin/bash
0.511  svstat          8747  8746  0   /command/svstat /service/nflx-httd
0.511  perl            8748  8746  0   /usr/bin/perl -e $l=<>;$l=~/(\d+) sec/;pr...
0.514  ps              8750  8749  0   /bin/ps --ppid 1 -o pid,cmd,args
0.514  grep            8751  8749  0   /bin/grep org.apache.catalina
0.514  sed              8752  8749  0   /bin/sed s/^ *//;
0.515  xargs           8754  8749  0   /usr/bin/xargs
0.515  cut              8753  8749  0   /usr/bin/cut -d   -f 1
0.523  echo             8755  8754  0   /bin/echo
0.524  mkdir            8756  8745  0   /bin/mkdir -v -p /data/tomcat
[...]
1.528  run            8785  1828  0   ./run
1.529  bash           8785  1828  0   /bin/bash
1.533  svstat          8787  8786  0   /command/svstat /service/nflx-httd
1.533  perl            8788  8786  0   /usr/bin/perl -e $l=<>;$l=~/(\d+) sec/;pr...
[...]
```

CPUs: runqlat

Scheduler latency (run queue latency)

```
# runqlat.py 10 1
Tracing run queue latency... Hit Ctrl-C to end.
```

usecs	:	count	distribution
0 -> 1	:	1906	***
2 -> 3	:	22087	*****
4 -> 7	:	21245	*****
8 -> 15	:	7333	*****
16 -> 31	:	4902	*****
32 -> 63	:	6002	*****
64 -> 127	:	7370	*****
128 -> 255	:	13001	*****
256 -> 511	:	4823	*****
512 -> 1023	:	1519	**
1024 -> 2047	:	3682	*****
2048 -> 4095	:	3170	****
4096 -> 8191	:	5759	*****
8192 -> 16383	:	14549	*****
16384 -> 32767	:	5589	*****

CPUs: runqlen

Run queue length

```
# runqlen.py 10 1
Sampling run queue length... Hit Ctrl-C to end.
```

runqlen	:	count	distribution
0	:	47284	*****
1	:	211	
2	:	28	
3	:	6	
4	:	4	
5	:	1	
6	:	1	

Memory: ffaults (book)

Page faults by filename

```
# ffaults.bt
Attaching 1 probe...
^C

[...]
@[dpkg]: 18
@[sudoers.so]: 19
@[ld.so.cache]: 27
@[libpthread-2.27.so]: 29
@[ld-2.27.so]: 32
@[locale-archive]: 34
@[system.journal]: 39
@[libstdc++.so.6.0.25]: 43
@[libapt-pkg.so.5.0.2]: 47
@[BrowserMetrics-5D8A6422-77F1.pma]: 86
@[libc-2.27.so]: 168
@[i915]: 409
@[pkgcache.bin]: 860
@[]: 25038
```

Disks: biolatency

Disk I/O latency histograms, per second

```
# biolatency.py -mT 1 5
Tracing block device I/O... Hit Ctrl-C to end.
```

06:20:16

msecs	: count	distribution
0 -> 1	: 36	*****
2 -> 3	: 1	*
4 -> 7	: 3	***
8 -> 15	: 17	*****
16 -> 31	: 33	*****
32 -> 63	: 7	*****
64 -> 127	: 6	*****

06:20:17

msecs	: count	distribution
0 -> 1	: 96	*****
2 -> 3	: 25	*****
4 -> 7	: 29	*****

[...]

File Systems: xfsslower

XFS I/O slower than a threshold (variants for ext4, btrfs, zfs)

```
# xfsslower.py 50
Tracing XFS operations slower than 50 ms

TIME      COMM          PID   T BYTES    OFF_KB    LAT(ms)  FILENAME
21:20:46  java          112789 R 8012     13925      60.16  file.out
21:20:47  java          112789 R 3571     4268       136.60  file.out
21:20:49  java          112789 R 5152     1780       63.88  file.out
21:20:52  java          112789 R 5214     12434      108.47  file.out
21:20:52  java          112789 R 7465     19379      58.09   file.out
21:20:54  java          112789 R 5326     12311      89.14   file.out
21:20:55  java          112789 R 4336     3051       67.89  file.out
[...]
22:02:39  java          112789 R 65536    1486748    182.10  shuffle_6_646_0.data
22:02:39  java          112789 R 65536    872492     30.10  shuffle_6_646_0.data
22:02:39  java          112789 R 65536    1113896    309.52  shuffle_6_646_0.data
22:02:39  java          112789 R 65536    1481020    400.31  shuffle_6_646_0.data
22:02:39  java          112789 R 65536    1415232    324.92  shuffle_6_646_0.data
22:02:39  java          112789 R 65536    1147912    119.37  shuffle_6_646_0.data
[...]
```

File Systems: xfsdist

XFS I/O latency histograms, by operation

```
# xfsdist.py 60
Tracing XFS operation latency... Hit Ctrl-C to end.
```

22:41:24:

```
operation = 'read'
```

usecs	: count	distribution
0 -> 1	: 382130	*****
2 -> 3	: 85717	*****
4 -> 7	: 23639	**
8 -> 15	: 5668	
16 -> 31	: 3594	
32 -> 63	: 21387	**

[...]

```
operation = 'write'
```

usecs	: count	distribution
0 -> 1	: 12925	****
2 -> 3	: 83375	*****

[...]

Networking: tcplife

TCP session lifespans with connection details

Networking: tcpsynbl (book)

TCP SYN backlogs as histograms

Languages: funccount

Count native function calls (C, C++, Go, etc.)

```
# funccount.py 'tcp_s*'  
Tracing 50 functions for "tcp_s*"... Hit Ctrl-C to end.  
^C  


| FUNC                          | COUNT  |
|-------------------------------|--------|
| [...]                         |        |
| tcp_setsockopt                | 1839   |
| tcp_shutdown                  | 2690   |
| tcp_sndbuf_expand             | 2862   |
| tcp_send_delayed_ack          | 9457   |
| tcp_set_state                 | 10425  |
| tcp_sync_mss                  | 12529  |
| tcp_sendmsg_locked            | 41012  |
| tcp_sendmsg                   | 41236  |
| tcp_send_mss                  | 42686  |
| tcp_small_queue_check.isra.29 | 45724  |
| tcp_schedule_loss_probe       | 64067  |
| tcp_send_ack                  | 66945  |
| tcp_stream_memory_free        | 178616 |
| Detaching...                  |        |


```

Applications: mysqlqld_qslower

MySQL queries slower than a threshold

```
# mysqlqld_qslower.py $(pgrep mysqld)
Tracing MySQL server queries for PID 9908 slower than 1 ms...
TIME(s)      PID      MS QUERY
0.000000    9962    169.032 SELECT * FROM words WHERE word REGEXP '^bre.*n$'
1.962227    9962    205.787 SELECT * FROM words WHERE word REGEXP '^bpf.tools$'
9.043242    9962    95.276 SELECT COUNT(*) FROM words
23.723025   9962    186.680 SELECT count(*) AS count FROM words WHERE word REGEXP
'bre.*n$'
30.343233   9962    181.494 SELECT * FROM words WHERE word REGEXP '^bre.*n$' ORDER BY word
[...]
```

Kernel: workq (book)

Work queue function execution times

Hypervisor: xenhyper (book)

Count hypercalls from Xen PV guests

```
# xenhyper.bt
Attaching 1 probe...
^C

@[mmu_update]: 44
@[update_va_mapping]: 78
@[mmuext_op]: 6473
@[stack_switch]: 23445
```

Containers: blkthrot (book)

Count block I/O throttles by blk cgroup

```
# blkthrot.bt
Attaching 3 probes...
Tracing block I/O throttles by cgroup. Ctrl-C to end
^C

@notthrottled[1]: 506


@throttled[1]: 31
```

That was only
14 out of
150+ tools

All are
open source

Not all 150+
tools shown here

Legend:
prior tool
new tool



Coping with so many BPF tools at Netflix

- On Netflix servers, **/apps/nflx-bpf-alltools** has all the tools
 - BCC, bpftrace, my book, Netflix internal
 - Open source at: <https://github.com/Netflix-Skunkworks/bpftoolkit>
- Latest tools are fetched & put in a hierarchy: cpu, disk, ...

```
bgregg@lgud-bgregg:~> ls --color ~/Git/nflx-bpf-alltools/root/apps/nflx-bpf-alltools/
applications/          disk/                  funcslower.py*  stackcount_example.txt
argdist_example.txt    filesystems/           hypervisors/   stackcount.py*
argdist.py*            funccount_example.txt  kernel/        tplist_example.txt
bpflist_example.txt    funccount.py*          languages/   tplist.py*
bpflist.py*          funclatency_example.txt  memory/       trace_example.txt
containers/             funclatency.py*        networking/  trace.py*
cpu/                   funcslower_example.txt  security/
```


- We are also building **GUIs** to front these tools

Tool development

Only one engineer at your company
needs to learn tool development

They can turn everyone's ideas into tools

The Tracing Landscape, Dec 2019



bcc/BPF (C & Python)

```
# load BPF program
b = BPF(text="""
#include <uapi/linux/ptrace.h>
#include <linux/blkdev.h>
BPF_HISTOGRAM(dist);
int kprobe__blk_account_io_completion(struct pt_regs *ctx,
    struct request *req)
{
    dist.increment(bpf_log2l(req->_data_len / 1024));
    return 0;
}
""")
```

```
# header
print("Tracing... Hit Ctrl-C to end.")

# trace until Ctrl-C
try:
    sleep(9999999)
except KeyboardInterrupt:
    print

# output
b["dist"].print_log2_hist("kbytes")
```

bcc examples/tracing/bitehist.py
entire program

bpftrace/BPF

```
bpftrace -e 'kr:vfs_read { @ = hist(retval); }'
```

<https://github.com/iovisor/bpftrace>
entire program

bpftrace Syntax

```
bpftrace -e 'k:do_nanosleep /pid > 100/ { @[comm]++ }'
```

Probe

Filter
(optional)

Action

Probe Type Shortcuts

tracepoint	t	Kernel static tracepoints
usdt	U	User-level statically defined tracing
kprobe	k	Kernel function tracing
kretprobe	kr	Kernel function returns
uprobe	u	User-level function tracing
uretprobe	ur	User-level function returns
profile	p	Timed sampling across all CPUs
interval	i	Interval output
software	s	Kernel software events
hardware	h	Processor hardware events

Filters

- /pid == 181/
- /comm != "sshd"/
- /@ts[tid]/

Actions

- Per-event output
 - `printf()`
 - `system()`
 - `join()`
 - `time()`
- Map summaries
 - `@ = count()` or `@++`
 - `@ = hist()`
 - ...

The following is in the https://github.com/iovisor/bpftrace/blob/master/docs/reference_guide.md

Functions

- **hist(n)** Log2 histogram
- **lhist(n, min, max, step)** Linear hist.
- **count()** Count events
- **sum(n)** Sum value
- **min(n)** Minimum value
- **max(n)** Maximum value
- **avg(n)** Average value
- **stats(n)** Statistics
- **str(s)** String
- **ksym(p)** Resolve kernel addr
- **usym(p)** Resolve user addr
- **kaddr(n)** Resolve kernel symbol
- **uaddr(n)** Resolve user symbol
- **printf(fmt, ...)** Print formatted
- **print(@x[, top[, div]])** Print map
- **delete(@x)** Delete map element
- **clear(@x)** Delete all keys/values
- **reg(n)** Register lookup
- **join(a)** Join string array
- **time(fmt)** Print formatted time
- **system(fmt)** Run shell command
- **cat(file)** Print file contents
- **exit()** Quit bptrace

Variable Types

- Basic Variables
 - **@global**
 - **@thread_local[tid]**
 - **\$scratch**
- Associative Arrays
 - **@array[key] = value**
- Builtins
 - **pid**
 - **...**

Builtin Variables

- **pid** Process ID (kernel tgid)
- **tid** Thread ID (kernel pid)
- **cgroup** Current Cgroup ID
- **uid** User ID
- **gid** Group ID
- **nsecs** Nanosecond timestamp
- **cpu** Processor ID
- **comm** Process name
- **kstack** Kernel stack trace
- **ustack** User stack trace
- **arg0, arg1, ...** Function args
- **retval** Return value
- **args** Tracepoint args
- **func** Function name
- **probe** Full probe name
- **curtask** Curr task_struct (u64)
- **rand** Random number (u32)

bpftrace: BPF observability front-end

```
# Files opened by process
bpftrace -e 't:syscalls:sys_enter_open { printf("%s %s\n", comm,
    str(args->filename)) }'

# Read size distribution by process
bpftrace -e 't:syscalls:sys_exit_read { @[comm] = hist(args->ret) }'

# Count VFS calls
bpftrace -e 'kprobe:vfs_* { @[func]++ }'

# Show vfs_read latency as a histogram
bpftrace -e 'k:vfs_read { @[tid] = nsecs }
    kr:vfs_read /@[tid]/ { @ns = hist(nsecs - @[tid]); delete(@tid) }'

# Trace user-level function
bpftrace -e 'uretprobe:bash:readline { printf("%s\n", str(retval)) }'
...
```

Example: bpftrace biolatency

Disk I/O latency histograms, per second

Example: bpftrace biolatency

Implemented in <20 lines of bpftrace

```
#!/usr/local/bin/bpftrace

BEGIN
{
    printf("Tracing block device I/O... Hit Ctrl-C to end.\n");
}

kprobe:blk_account_io_start
{
    @start[arg0] = nsecs;
}

kprobe:blk_account_io_done
/@start[arg0]/
{
    @usecs = hist((nsecs - @start[arg0]) / 1000);
    delete(@start[arg0]);
}
```

BCC/BPF: biolatency

100.66.98.191:7402



?88-1048575

?144-524287

!072-262143

!536-131071

!2768-65535

!6384-32767

8192-16383

4096-8191

2048-4095

1024-2047

512-1023

256-511

128-255

64-127

32-63


16-31

8-15

4-7

2-3

0-1



11:26:20

11:26:40


11:27:00

11:27:20

11:27:40

11:28:00

Netflix Vector
(old)




Grafana at Netflix

Takeaways

Add **BCC & bpftrace** packages to your servers

Start using BPF perf tools directly or via GUIs

Identify 1+ engineer at your company to develop tools & GUIs



From: BPF Performance Tools: Linux System and Application Observability, Brendan Gregg, Addison Wesley 2019

Thanks & URLs



BPF: Alexei Starovoitov, Daniel Borkmann, David S. Miller, Linus Torvalds, BPF community

BCC: Brenden Blanco, Yonghong Song, Sasha Goldstein, BCC community

bpftrace: Alastair Robertson, Matheus Marchini, Dan Xu, bpftrace community


<https://github.com/iovisor/bcc>

<https://github.com/iovisor/bpftrace>

<https://github.com/brendangregg/bpf-perf-tools-book>

<http://www.brendangregg.com/ebpf.html>

<http://www.brendangregg.com/bpf-performance-tools-book.html>



All diagrams and photos (slides 11 & 22) are my own; slide 12 is from KernelRecipes: <https://www.youtube.com/watch?v=bbHFg9lsTk8>

Thank you!

Brendan Gregg

@brendangregg
bgregg@netflix.com



Please complete the session
survey in the mobile app.